

Statistical Evidence for a DQN-Family Advantage in a Resource-Coupled Satellite-Ground Scheduling Benchmark

Bocheng Yang ¹, Lujie Zheng ¹, Dingyang Lyu ³, Jian Liu ^{1,2,*} and Bo Chen ^{1,*}

¹ School of Aerospace, Harbin Institute of Technology, Shenzhen 518055, China; 25b961001@stu.hit.edu.cn (B.Y.); 24s061021@stu.hit.edu.cn (L.Z.);

² Guangxi Key Laboratory of Precision Navigation Technology and Application, Guilin University of Electronic Technology, Guilin 541004, China;

³ Zhongguancun Academy, Beijing 100094, China; s-ldy25@bza.edu.cn (D.L.)

* Correspondence: hitchenbo@hit.edu.cn (B.C.); liujian2024@hit.edu.cn (J.L.)

Highlights

What are the main findings?

- All 30 paired DQN–MPC-4 objective–regime comparisons supported a DQN-family advantage.
- Within-family differences were smaller and scenario dependent, while deeper exact planning reduced cost but required substantially more decision time.

What are the implications of the main findings?

- Long-horizon value learning can outperform short-horizon planning in this synthetic, soft-constrained, resource-coupled satellite-ground scheduling benchmark.
- The reproducible benchmark provides a foundation for mission-calibrated, hard-constrained, and operational validation.

Abstract

Resource-coupled satellite-ground scheduling requires decisions with computational, energy, thermal, and communication effects that persist across task arrivals. We tested whether the deep Q-network (DQN) family provides a reproducible advantage over short-horizon planning in this setting. Six DQN objectives used the same state representation, network architecture, replay procedure, optimizer, interaction budget, 20 model seeds, and held-out traces. We evaluated them across five fully coupled regimes in a fully specified synthetic benchmark with soft constraints. In all 30 objective–regime comparisons, the mean episode cost for each DQN objective was lower than that of four-step model-predictive control (MPC-4). Every paired 95% bootstrap interval excluded zero, and all Holm-adjusted sign tests supported the same ordering. The least-favorable mean difference was -5.92, and the family-wide simultaneous 95% upper bound was -5.03. Differences among DQN objectives were smaller and scenario dependent. Extending exact planning from one to six steps reduced cost but substantially increased decision time. These results support a DQN-family advantage across objective designs and resource-stress regimes within this synthetic benchmark. The reproducible platform provides a basis for mission-calibrated, hard-constrained, and operational validation.

Keywords: deep Q-network variants; satellite edge computing; task scheduling; model-assisted reinforcement learning; full-action supervision

Received:

Revised:

Accepted:

Published:

Copyright: © 2026 by the authors.

Submitted to *Remote Sens.* for possible open access publication under the

terms and conditions of the [Creative](#)

[Commons Attribution \(CC BY\)](#) license.




Figure 1. Family-level evidence for deep Q-network (DQN) methods relative to four-step model-predictive control (MPC-4) in the synthetic benchmark with soft constraints. (a) The evaluation uses 20 paired model-seed blocks, each averaging 20 shared held-out traces. (b) Cells report the mean episode-cost difference between each DQN objective and MPC-4; negative values favor DQN. (c) Circles and lines show the largest, least-favorable mean difference within each scope and its 95% joint-bootstrap interval. Open triangles show simultaneous 95% upper bounds. A single 10,000-resample seed-block bootstrap preserves correlation among the 30 contrasts. Exact two-sided sign tests are Holm-adjusted across all comparisons, which satisfy the prespecified support rule.

1. Introduction

Earth-observation and communication satellites increasingly execute perception and decision workloads in orbit. Compact onboard networks can reduce downlink demand and response latency under constrained computational resources [1–5]. Onboard execution consumes heat, energy, and computing capacity, whereas ground offloading consumes contact time and bandwidth. Scheduling must therefore coordinate onboard execution, complete offloading, and collaborative preprocessing over time.

Existing schedulers optimize constellation coordination, observation plans, service placement, deadlines, makespan, or energy use [6–12]. In our setting, every decision changes six resources that persist across task arrivals: heat, energy, queue occupancy, compute utilization, bandwidth, and remaining contact time. An action with a low immediate cost can therefore restrict the choices available for later tasks. This temporal coupling raises a central question: can learned long-horizon values outperform transparent one-step and finite-horizon alternatives across changing resource conditions?

Standard deep Q-network (DQN) learning [13,14] constructs one Bellman target from each factual tuple (s_t, a_t, r_t, s_{t+1}) . In our setting, analytical resource equations can also evaluate every unexecuted action without advancing the environment or revealing later tasks. This structure enables controlled evaluation of the DQN family and alternative uses of full-action supervision during training.

We compare six DQN objectives: standard DQN, full-action Q, immediate advantage, centered full-action, Double DQN, and Double DQN with centered full-action supervision. A separate $\gamma = 0$ contextual bandit provides a learned one-step comparator. Model-evaluated alternatives are used only during training. During evaluation and deployment, every policy acts directly from its learned Q-values.

We make three contributions:

1. We introduce a fully specified, reproducible benchmark that links task primitives to three processing pathways and six persistent computational, energy, thermal, and communication resources.
2. We construct a controlled six-objective DQN family that varies factual supervision, full-action targets, centering, bootstrapping, and Double DQN evaluation under matched training conditions.
3. We test the family-level result against four-step model-predictive control (MPC-4) through 30 paired objective–regime contrasts, a correlation-preserving joint bootstrap, multiplicity-corrected tests, and a simultaneous upper bound.

2. Related Work and Study Positioning

2.1. Satellite-ground scheduling

Satellite schedulers address constellation coordination, observation planning, service placement, and energy-aware offloading [6–12]. Our setting combines an explicit cost for every action with resource carryover between task arrivals. Because each immediate action cost is directly evaluable, the one-step objective is observable. The benchmark therefore asks whether a learned value function can exploit consequences that persist beyond this immediate cost.

2.2. Model use and full-action learning

Dyna introduced simulated experience for value learning and planning [15]. Fitted dynamics can likewise accelerate deep Q-learning [16]. Counterfactual data fusion and causal augmentation construct transitions under structural assumptions [17,18]. Counterfactual credit assignment instead separates action influence from future randomness [19]. Recent work has also examined multi-action generative interfaces and joint quantities across action outcomes [20]. Together, these studies provide several routes for incorporating modeled alternatives into reinforcement learning.

Advantage learning modifies Bellman operators to widen action gaps [21]. Four objectives in our DQN family instead retain their respective Bellman target rules while adding modeled outcomes for all actions at a sampled state. Table 1 positions this training-time supervision interface relative to related mechanisms. Standard DQN and Double DQN serve as factual-only family controls. We compare the DQN family with short-horizon planning and examine how target construction, centering, and target evaluation affect within-family variation.

In this manuscript, “counterfactual” denotes a model-evaluated one-step outcome for an unexecuted action under the same current task and exogenous state. It does not denote an identified causal effect from observational data.

Table 1. Positioning of training-time modeled-action supervision used by four deep Q-network (DQN) objectives. “All actions” means that one update can use modeled outcomes for every discrete action at the sampled state. Standard DQN and Double DQN provide factual-only family controls.

Method family	Source of modeled information	How it enters learning	Relation to the evaluated supervision interface
Dyna and local model-based acceleration [15,16]	Learned or fitted dynamics	Simulated experience or local rollouts	Does not use one simultaneous analytical outcome per action in a training-only auxiliary objective
Causal counterfactual augmentation [18]	Structural assumptions and offline trajectories	Synthetic transitions augment the training distribution	Targets distributional coverage rather than within-state Bellman action gaps
Counterfactual credit assignment [19]	Future-conditioned trajectory information	Policy-gradient baseline or critic	Model-free and trajectory-conditioned rather than analytical one-step action enumeration
Joint Markov decision processes (MDPs) [20]	Multi-action sample transition model	Joint action-outcome quantities and Bellman operators	Closest formal interface; not a training-only auxiliary objective on a fixed factual DQN backbone
Evaluated full-action DQN objectives	Deterministic analytical one-step transition for three actions	Auxiliary full-action targets or centered action gaps; no replay insertion	One controlled factor within a six-objective family; the family claim also includes factual-only controls

3. Materials and Methods

3.1. System Model and Immediate-Cost Quantification

At each task arrival t , a satellite selects one of three actions. Action α_s executes the task onboard and downlinks the result, whereas α_g downlinks the complete input for ground execution. Action $\alpha_{s\&g}$ performs onboard preprocessing followed by ground execution:

$$\mathcal{A} = \{\alpha_s, \alpha_g, \alpha_{s\&g}\},$$

(1)

where the subscripts denote satellite, ground, and collaborative processing, respectively.




Figure 2. Three satellite-ground processing pathways. Each arriving task is executed onboard, transferred for ground execution, or preprocessed collaboratively before ground completion. The analytical transition model returns the immediate cost and the post-action heat, energy, queue occupancy, compute utilization, bandwidth, and remaining contact.

3.1.1. Implemented immediate-cost model

The experiments use the dimensionless linear cost implemented in the internally consistent synthetic simulator. Let x_k denote the 15 task descriptors in Table 2. Each descriptor is normalized as

$$\bar{x}_k = \text{clip}(x_k/s_k, 0, 1.5), \quad (2)$$

where s_k is the corresponding scale. For inherited normalized bandwidth b , the available return rate is $R(b) = 2.2 + 217.8b \text{ Mbit s}^{-1}$. The three action latencies are

$$\begin{aligned} L_s &= x_3 + 1000x_6/R(b), \\ L_g &= x_8 + 1000x_9/R(b), \\ L_{s\&g} &= 1000x_{12}/4 + x_{13} + 1000x_{15}/R(b). \end{aligned} \quad (3)$$

The implemented immediate costs are

$$\begin{aligned} c_s^{imm} &= 0.70\bar{x}_1 + 0.65\bar{x}_2 + 0.00045L_s + 0.55\bar{x}_5 + 0.35\bar{x}_6 + 0.35h + 0.12(1 - e), \\ c_g^{imm} &= 0.45\bar{x}_7 + 0.00045L_g + 0.55\bar{x}_9 + 0.35(0.28 - \tau)_+, \\ c_{s\&g}^{imm} &= 0.55(\bar{x}_{10} + \bar{x}_{11}) + 0.55\bar{x}_{14} + 0.00045L_{s\&g} + 0.45\bar{x}_{15} \\ &\quad + 0.18h + 0.12q + 0.18(0.22 - \tau)_+. \end{aligned} \quad (4)$$

Here h , e , q , and τ denote inherited heat, remaining energy, queue occupancy, and contact duration, respectively. The weights are fixed simulator assumptions rather than fitted mission utilities. The same weights are used during training, counterfactual preview, MPC evaluation, and deployment evaluation. Automated tests reconstruct Eq. (4) from the 15 raw descriptors and compare the resulting costs with all three implementation outputs.

Table 2. Implemented 15-dimensional task descriptor dictionary. The complete 24-dimensional observation adds normalized time, the sine and cosine of link phase, and six resource coordinates: h , e , q , compute utilization u , bandwidth b , and contact duration τ .

Coordinate	Descriptor	Unit	Normalization scale s_k
x_1	Onboard compute heat	J	100
x_2	Total workload	GFLOP	4.8
x_3	Onboard compute time	ms	1400
x_4	Result transmit time at 50 Mbit s ⁻¹	ms	100
x_5	Squared co-location count	count ²	2401
x_6	Onboard-result data	Mbit	4
x_7	Full-offload transmit heat	J	100
x_8	Ground compute time	ms	100
x_9	Raw input data	Mbit	76.8
x_{10}	Preprocessing heat	J	100
x_{11}	Feature-transmit heat	J	100
x_{12}	Onboard preprocessing work	GFLOP	2.2
x_{13}	Remaining ground compute time	ms	100
x_{14}	Total workload descriptor	GFLOP	4.8
x_{15}	Collaborative feature data	Mbit	32

3.1.2. Internally consistent task and communication generator

Rather than sampling heat, workload, data volume, and latency independently, the simulator first draws correlated latent variables (u_D, u_W) with correlation coefficient 0.68. A logistic transform maps these variables to raw data volume $D_t \in [8, 64]$ Mbit and workload $W_t \in [0.25, 4.0]$ GFLOP. Result and collaborative-feature volumes are

$$D_t^{res} = r_t D_t, \quad r_t \in [0.01, 0.05], \quad D_t^{feat} = f_t D_t, \quad f_t \in [0.08, 0.30]. \quad (5)$$

The implementation enforces $D_t^{res} \leq D_t^{feat} \leq D_t$. A workload-dependent fraction $p_t \in [0.20, 0.45]$ is processed onboard in the collaborative mode. Given onboard and ground rates $R_s = 4$ and $R_g = 80$ GFLOP s⁻¹, processing time is computed as W/R rather than sampled separately. Compute power increases monotonically from 4 to 20 W with workload, and compute heat equals 0.78 times compute energy. Transmission time is the data volume divided by the effective return rate


$$R_t^{link} = 2.2 + 217.8b_t \quad \text{Mbit s}^{-1}, \quad (6)$$

and transmission heat equals 0.55 times transmit energy. These values define engineering envelopes rather than fitted flight distributions. Their scales are consistent with modern SmallSat computing-system power reported by the National Aeronautics and Space Administration (NASA). NASA ground-network examples range from 2.2 Mbit s⁻¹ S-band to 220 Mbit s⁻¹ X-band returns [22,23]. The bounded data-reduction factors model the bandwidth-reduction role of onboard image compression rather than any instrument-specific codec [24]. Table 3 lists every implemented assumption.

Table 3. Modeled engineering primitives and provenance of the current simulator. Values define a transparent synthetic envelope; they are not presented as a fitted flight distribution.

Quantity	Implemented range	Basis
Raw task data	8–64 Mbit nominal	Stress interval reaches 92.8 Mbit before sensitivity scaling
Arithmetic workload	0.25–4.0 GFLOP nominal	Stress interval reaches 5.8 GFLOP
Result / raw-data ratio	1–5%	Bounded compression assumption
Feature / raw-data ratio	8–30%	Bounded collaborative-processing assumption
Preprocessing fraction	20–45%	Deterministic function of workload quantile
Onboard / ground rate	4 / 80 GFLOP s ^{−1}	Fixed simulator hardware envelope
Compute power	4–20 W	Bracketed by NASA SmallSat avionics examples
Effective RF return rate	2.2–220 Mbit s ^{−1}	NASA SmallSat ground-system examples
Transmit power	6 W	Fixed simulator assumption
Burst multiplier	1.45×	Contiguous prespecified stressor

Before training, an automated audit generated 20,000 tasks. All values were finite and positive, all data-flow inequalities held, and the empirical raw-data–workload correlation was 0.660. Figure 3 shows the resulting dependence structure. Constant-rate columns have zero variance and are therefore excluded from the correlation analysis.

**Figure 3.** Internal-consistency audit of the synthetic generator. (a) A legibility sample of 2,000 of 20,000 tasks; color encodes heat derived from compute energy. (b) Correlations computed from all 20,000 tasks. No fitted flight distribution is implied.

3.1.3. Complete stochastic generator and action-load specification

To support independent reconstruction, this subsection specifies the remaining simulator functions used in the transition. Let $u_D, \epsilon_W \sim \mathcal{N}(0,1)$ independently, $u_W = 0.68u_D + \sqrt{1 - 0.68^2}\epsilon_W$, and $g(u) = [1 + \exp(-1.55u)]^{-1}$. Nominal raw volume and workload are

$$D = 8 + 56g(u_D) \text{ Mbit}, \quad W = 0.25 + 3.75g(u_W) \text{ GFLOP}. \quad (7)$$

The burst interval spans indices $\lfloor T/3 \rfloor$ through $\lfloor T/3 \rfloor + \max(4, \lfloor T/4 \rfloor) - 1$ and multiplies both quantities by 1.45. With independent $U_r, U_f \sim \mathcal{U}(0,1)$,

$$\begin{aligned} r &= 0.01 + 0.04U_r, & f &= 0.08 + 0.22 \text{clip}[0.65g(u_W) + 0.35U_f, 0, 1], \\ p &= 0.20 + 0.25g(u_W), & D^{\text{res}} &= rD, \quad D^{\text{feat}} = fD, \quad W^{\text{pre}} = pW, \\ n &= \text{round}\{6 + 43 \text{clip}[0.60g(u_W) + 0.40g(u_D), 0, 1]\}, & x_5 &= n^2. \end{aligned} \quad (8)$$

Onboard and ground rates are 4 and 80 GFLOP s^{−1}, respectively. Compute power is $P_c = 4 + 16 \text{clip}(W/4.8, 0, 1)^{0.70}$ W, and transmit power is 6 W. Compute and transmit heat

equal 0.78 and 0.55 times their corresponding energy use, respectively. These definitions generate all 15 coordinates in Table 2.

At reset, a single phase $\phi \sim \mathcal{U}(0, 2\pi)$ generates the exogenous link trace

$$\begin{aligned}\bar{b}_t &= \text{clip}\left\{\kappa_b[0.62 + 0.23 \sin(2\pi t/24 + \phi) + \epsilon_t^b] - I_{link}, 0.12, 1\right\}, \\ \bar{\tau}_t &= \text{clip}\{0.55 + 0.35 \sin(2\pi t/31 + \phi/2) + \epsilon_t^\tau, 0.08, 1\},\end{aligned}\quad (9)$$

where $\epsilon_t^b \sim \mathcal{N}(0, 0.06^2)$, $\epsilon_t^\tau \sim \mathcal{N}(0, 0.04^2)$, κ_b is the sensitivity multiplier, and $I_{link} = 0.22$ only in the link-limited regime. All tasks and link states for an episode are drawn at reset with one seeded NumPy generator.

The normalized action loads used in Eq. (12) are stated here in action order $(\alpha_s, \alpha_g, \alpha_{s\&g})$:

$$\begin{aligned}\hat{w} &= (W, 0.05W, W)/4.8, & \hat{Q} &= (x_1, 0.20x_7, x_{10} + x_{11})/100, \\ \hat{B} &= (x_6, x_9, x_{15})/76.8, \\ \Delta e &= \kappa_e(0.012 + 0.042\hat{w}_1, 0.007 + 0.025\hat{B}_2/b, 0.010 + 0.025\hat{w}_3 + 0.018\hat{B}_3/b), \\ \Delta q &= (0.055\hat{w}_1(1 + u), 0.048\hat{B}_2/b, 0.030(\hat{w}_3 + \hat{B}_3/b)), \\ \mu &= (0.030(1 - u), 0.035b\tau, 0.032(1 - 0.5u)), \\ \ell &= (0.02, 0.18\hat{B}_2, 0.12\hat{B}_3).\end{aligned}\quad (10)$$

Each division by b uses $\max(b, 0.1)$. The utilization targets in Eq. (12) are $(\hat{w}_1, 0.10\hat{B}_2, 0.55\hat{w}_3)$. Initial resources are $(h, e, q, u) = (0.22, 0.88, 0.12, 0.18)$. Energy-limited episodes instead use $e = 0.62$, and thermal-stress episodes use $h = 0.42$. Each episode terminates after exactly 64 executed actions, and the terminal Bellman multiplier is zero. Energy and the other resources remain clipped, soft-constrained accounting states. In particular, $e = 0$ neither masks infeasible actions nor terminates an episode. It therefore does not represent a conserved battery state, and clipping does not retain the magnitude of an energy deficit.

3.2. Sequential Scheduling Formulation

3.2.1. State and transition

The normalized 24-dimensional state is

$$s_t = [x_t, \eta_t, z_t], \quad \eta_t = [t/(T - 1), \sin \phi, \cos \phi], \quad z_t = [h_t, e_t, q_t, u_t, b_t, \tau_t], \quad (11)$$

where $x_t \in \mathbb{R}^{15}$ contains the task inputs defined in Table 2. The context vector η_t exposes absolute episode time and the reset-level link phase. Together, these variables determine the time-varying task and link distributions. The remaining task and link innovations are independent draws conditional on this context. The resource vector contains normalized heat, remaining energy, queue occupancy, compute utilization, bandwidth, and contact duration. Thus, within each fixed regime, the transition distribution is conditioned on all persistent simulator variables rather than on a hidden reset phase.

Executing action a_t changes resources before the next exogenous task arrives. The implemented transition is

$$\begin{aligned} h_{t+1} &= \text{clip}(0.86h_t + 0.17c\widehat{Q}_t^{(a_t)}, 0, 1.3), \\ e_{t+1} &= \text{clip}(e_t + 0.006 - c\Delta e_t^{(a_t)}, 0, 1), \\ q_{t+1} &= \text{clip}(q_t + c(\Delta q_t^{(a_t)} - \mu_t^{(a_t)}), 0, 1.3), \\ u_{t+1} &= \text{clip}(0.68u_t + 0.32c\widehat{w}_t^{(a_t)}, 0, 1.2), \\ b_{t+1} &= \text{clip}(\bar{b}_{t+1} - c\ell_t^{(a_t)}, 0.08, 1), \\ \tau_{t+1} &= \text{clip}(\bar{\tau}_{t+1} - 0.10c\widehat{B}_t^{(a_t)}, 0.05, 1). \end{aligned} \quad (12)$$

Here $\widehat{Q}$, $\widehat{w}$, and $\widehat{B}$ denote normalized action-specific heat, work, and data demands, respectively. The terms Δe , Δq , μ , and ℓ are the deterministic resource-use, queue, service, and link-occupation functions defined in Eq. (10). Exogenous link variables $\bar{b}$ and $\bar{\tau}$ follow Eq. (9) and are sampled before each episode. Setting $c = 0$ removes action-dependent carryover, $c = 0.5$ weakens it, and $c = 1$ produces fully coupled dynamics.

3.2.2. Reward and standard DQN backbone

The reward is the negative immediate cost plus post-decision feasibility penalties:

$$\begin{aligned} r_t = & -c_t^{imm} - c[12(h_{t+1} - 0.78)_+ + 15(0.18 - e_{t+1})_+ \\ & + 8(q_{t+1} - 0.75)_+ + 10\mathbf{1}_{a_t \neq a_s}(0.14 - \tau_t)_+]. \end{aligned} \quad (13)$$

The DQN uses the unmodified factual target

$$y_t = r_t + \gamma(1 - d_t) \max_{a'} Q_{\bar{\theta}}(s_{t+1}, a'), \quad (14)$$

and factual loss

$$L_{DQN} = \mathcal{H}(Q_{\theta}(s_t, a_t), y_t), \quad (15)$$

where $\mathcal{H}$ denotes the Huber loss. The target network both selects and evaluates the maximizing action, as required by standard DQN.

3.3. DQN Objective Family and Full-Action Supervision

All six objectives share the factual Huber loss and the same implementation settings. Standard DQN members use Eq. (14), whereas Double DQN members use online-network action selection with target-network evaluation. Standard DQN and Double DQN learn only from factual transitions. Four model-assisted objectives additionally use a training-only auxiliary path that supplies labels for unexecuted actions. The following construction defines the shared preview interface and the centered member. The experiments also evaluate uncentered full-action Q, immediate advantage, and the corresponding Double DQN auxiliary objective.

3.3.1. Side-effect-free full-action preview

For every replayed factual state, the analytical transition evaluates each action $a \in \mathcal{A}$:

$$(\tilde{\ell}_{t,a}, \tilde{z}'_{t,a}) = F_{preview}(x_t, z_t, a). \quad (16)$$

Here $\tilde{\ell}_{t,a}$ contains the immediate cost and post-decision feasibility penalty for action a . The preview does not advance time, mutate state, alter trajectories, or consume random numbers. Before action selection, it computes only the current-task cost and six post-action

resources. The factual action is then executed once. After the next exogenous task x_{t+1} is revealed, the replay tuple is completed. During a later replay update, all alternatives reuse this observed next task and exogenous link trace. The counterfactual next state therefore replaces only the six endogenous resource coordinates:

$$\tilde{s}'_{t,a} = [x_{t+1}, \eta_{t+1}, \tilde{z}'_{t,a}]. \quad (17)$$

Evaluation and deployment never call the preview interface.

3.3.2. All-action counterfactual Bellman targets

The simulator-exact one-step outcome yields

$$\tilde{y}_{t,a} = -\tilde{\ell}_{t,a} + \gamma(1 - d_t) \max_{a'} Q_{\tilde{\theta}}(\tilde{s}'_{t,a}, a'). \quad (18)$$

These targets require no multi-step rollout, and no counterfactual transition is inserted into replay. The centered full-action objective uses the targets only after the state-wise centering defined below.

3.3.3. Centered advantage distillation

Absolute targets in Eq. (18) share a bootstrapped state-value offset. Deployment selects arg max_a $Q(s, a)$, so a common additive offset cannot change the action. We define predicted and target centered advantages as

$$\begin{aligned} \hat{A}_{t,a} &= Q_{\theta}(s_t, a) - \frac{1}{|\mathcal{A}|} \sum_j Q_{\theta}(s_t, j), \\ \tilde{A}_{t,a} &= \tilde{y}_{t,a} - \frac{1}{|\mathcal{A}|} \sum_j \tilde{y}_{t,j}. \end{aligned} \quad (19)$$

The advantage and total losses are

$$\begin{aligned} L_{adv} &= \frac{1}{|\mathcal{A}|} \sum_{a \in \mathcal{A}} \mathcal{H}(\hat{A}_{t,a}, \tilde{A}_{t,a}), \\ L_{CFA} &= L_{DQN} + \lambda L_{adv}. \end{aligned} \quad (20)$$

The factual term anchors the absolute Q scale, while L_{adv} teaches all action gaps. The only added algorithmic hyperparameter is λ .

Algorithm 1: DQN training with centered full-action supervision

1. Initialize uniform replay $\mathcal{D}$, online Q_{θ} , and target $Q_{\tilde{\theta}} \leftarrow Q_{\theta}$.
2. At each real step, observe s_t and preview current-task outcomes $(\tilde{\ell}_{t,a}, \tilde{z}'_{t,a})$ for all actions without side effects.
3. Select a_t by the shared epsilon-greedy schedule and execute it once. After x_{t+1} is revealed, store the factual tuple plus the previously computed previews.
4. When an update is due, sample one uniform mini-batch, compute the factual loss, compute Eqs. (18) and (19), and update θ using Eq. (20).
5. Synchronize $Q_{\tilde{\theta}}$ at the shared fixed interval.

Training previews are never used for evaluation or deployment.

3.4. Experimental Design and Settings

3.4.1. Design principles

The design separates primary family-level inference from secondary within-family analyses. The primary analysis compares all six DQN objectives with MPC-4 across five fully coupled regimes. This design yields 30 objective–regime contrasts under shared seed blocks and traces. Secondary analyses examine how auxiliary supervision and Double DQN target evaluation affect performance within the family. Comparing standard DQN with its centered full-action variant isolates the addition of that auxiliary objective to an otherwise unchanged factual DQN backbone. Full-action Q and immediate advantage use the same three modeled outcomes per state, matching the type and number of model outputs. However, these comparisons do not isolate centering or Bellman bootstrapping because target scales differ and auxiliary weights were not tuned separately. Comparing Double DQN with and without centered full-action supervision assesses the descriptive effect when action selection and target evaluation are decoupled [25]. The $c \in \{0, 0.5, 1\}$ regimes and $H \in \{1, 2, 4, 6\}$ planning horizons bound the comparison across coupling and planning depth.

Standard DQN and its centered full-action variant differ only in the auxiliary objective. They share the same input, network, replay, exploration, optimizer, update cadence, target synchronization, and gradient clipping. The full-action Q control applies the same auxiliary weight directly to the three uncentered Bellman targets. The immediate-advantage control centers the negative analytical one-step costs but omits their bootstrapped future values. Both Double DQN variants use online-network action selection with target-network evaluation in factual and auxiliary Bellman targets. Dueling networks, prioritized replay, multi-step returns, noisy networks, and distributional value learning remain outside the comparison.

We fixed the auxiliary weight before running experiments with the internally consistent simulator. A preceding calibration with model seeds 700–703 compared $\lambda \in \{0.1, 0.3, 1, 3\}$. The rule selected $\lambda = 0.3$ by maximizing the worst relative improvement of the centered full-action variant across five coupled regimes. Ties favored the smaller weight. This rule was specific to the centered objective. The uncentered full-action, immediate-advantage, and Double DQN auxiliary variants inherited $\lambda = 0.3$ without method-specific calibration. Because the auxiliary losses differ in scale and error structure, these are fixed-coefficient implementation comparisons. They do not factorially isolate centering or bootstrapping. The calibration used an earlier simulator version and therefore serves as a prior specification rather than evidence from the current generator. Confirmatory and component analyses use seeds 800–819 without post hoc removal, replacement, or selection. Every learned policy was trained only on the nominal $c = 1$ distribution.

3.4.2. Training configuration

Each network has 24 inputs, hidden widths of 128–128–64, rectified linear unit (ReLU) activations, and three outputs. Every model receives 600 episodes of 64 tasks, totaling 38,400 real interactions. Fixed settings are $\gamma = 0.97$, a learning rate of 3×10^{-4} , a batch size of 128, a replay capacity of 100,000, and a gradient-norm limit of 10. Training begins after 128 transitions and performs one update every four interactions. The checkpoint field `target_interval=250` is evaluated only on update calls. Because updates occur at interaction counts divisible by four, synchronization occurs every 500 real interactions, equivalent to 125 optimizer updates. Epsilon decreases linearly from 1 to 0.05 over 13,440 interactions.

The contextual bandit uses the same network, optimizer, exploration schedule, and interaction budget, with $\gamma = 0$. It therefore learns the modeled one-step reward without a

long-horizon Bellman term. Uniform replay stores only one-step factual transitions, with preview outputs attached as auxiliary training data. Standard DQN and Double DQN do not use analytical labels for alternative actions. The four auxiliary variants evaluate all three modeled action outcomes at every real interaction. The checkpoint manifest records these counts, optimizer updates, and wall time for every trained model.

All reported experiments use the internally consistent synthetic generator described above. Link bandwidth and contact duration follow bounded, noisy periodic traces generated at reset. Burst and thermal-stress episodes multiply raw data and workload by 1.45 during one contiguous interval. All dependent heat, energy, latency, and transmission quantities are then recomputed from those primitives. An earlier independently sampled generator is retained only as a historical reproducibility mode and is not used for the reported results.

3.4.3. Seven evaluation regimes

The seven regimes comprise two composite coupling-boundary conditions and five fully coupled operating conditions:

1. **Static** ($c = 0$): removes action-dependent state transitions and scales delayed feasibility penalties to zero.
2. **Reduced coupling** ($c = 0.5$): halves action-dependent resource effects and the delayed feasibility-penalty scale.
3. **Nominal** ($c = 1$): matches the fully coupled training distribution.
4. **Burst**: increases contiguous compute and thermal demand.
5. **Link-limited**: reduces the available bandwidth trace by 0.22 before clipping.
6. **Energy-limited**: lowers initial energy from 0.88 to 0.62.
7. **Thermal-stress**: raises initial heat from 0.22 to 0.42 and includes the burst interval.

The first two regimes are composite interventions because c scales both the transition in Eq. (12) and the delayed-penalty term in Eq. (13). They therefore do not isolate action-dependent carryover from reward rescaling. The remaining five regimes retain $c = 1$. Models trained at $c = 1$ are transferred to the two boundary conditions without retraining. These rows measure behavior under a joint objective-and-dynamics shift rather than the best policy attainable at lower coupling.

3.4.4. Planning comparator, outcomes, and inference

The deployment comparison includes immediate-cost argmin, a contextual bandit, a prespecified resource-threshold heuristic, standard DQN, the centered full-action family member, and an exact H -step receding-horizon planner. The threshold heuristic first checks whether $h > 0.72$ or $e < 0.24$. If this condition holds and both $b > 0.28$ and $\tau > 0.18$, it selects ground processing; otherwise, it selects collaborative processing. When neither heat nor energy triggers the rule, it selects onboard processing if $b < 0.22$ or $\tau < 0.14$. If this bandwidth or contact condition is also absent, it selects collaborative processing when $q > 0.70$ and immediate-cost argmin otherwise. These thresholds were specified manually and were not fitted to the evaluation traces. Fixed onboard and fixed ground policies always select α_s and α_g , respectively. At each decision, the planner enumerates all 3^H action sequences over the next H known tasks and link states. It minimizes discounted modeled cost with $\gamma = 0.97$, executes only the first action, and replans. The primary table reports $H = 4$. A separate nominal analysis evaluates $H \in \{1, 2, 4, 6\}$ on 400 unique held-out traces and reports wall-clock time per decision. Each planner is exact only for its stated perfect-forecast objective, not for the 64-step global problem.

The primary outcome is the undiscounted episode total cost

$$C_{\text{episode}} = \sum_{t=0}^{T-1} (c_t^{\text{imm}} + p_t) = - \sum_{t=0}^{T-1} r_t, \quad (21)$$

where p_t is the complete nonnegative penalty inside the square brackets in Eq. (13), including its multiplier c . Every executed step, including the terminal step, contributes once. DQN and the finite-horizon planner optimize discounted targets with $\gamma = 0.97$. In contrast, Eq. (21) reports accumulated operational cost without discounting. We did not conduct a separate sensitivity analysis of this optimization–evaluation distinction. Secondary outcomes are immediate cost, delayed penalty, energy use, modeled latency, transmitted data, threshold-exposure steps, planner wall-clock time, maximum heat, and minimum remaining energy.

The independent inferential unit is the model-seed block. For learned–learned comparisons, the same seed initializes network weights and controls task seeds, exploration draws, and replay-index sampling. Trajectories may diverge after different greedy actions, but the matched seed remains the randomization block. Each seed is evaluated on the same 20 held-out traces, which are averaged within seed. This procedure gives $n = 20$ blocks per learned method. For the central DQN-family–MPC-4 claim, each raw row is matched by regime, model-seed trace namespace, and trace seed. Trace-level differences are then averaged within seed. MPC-4 is deterministic conditional on a trace; its model-seed field therefore indexes a matched trace namespace rather than a trained planner replicate.

All paired bootstrap analyses use 10,000 resamples with seed 20260829. Individual two-sided 95% intervals are percentile intervals, and exact two-sided paired sign tests assess directional consistency. For the 30 DQN-objective–regime contrasts against MPC-4, one joint bootstrap resamples the same 20 seed blocks across all contrasts. This procedure preserves correlation among the contrasts. Sign-test values are Holm-adjusted across all cells. We also report the largest, least-favorable mean difference within each regime and across the complete contrast set. A simultaneous one-sided 95% upper bound adds a joint critical value to every observed cell mean. The critical value is the 95th percentile of the largest centered bootstrap deviation across the relevant cells. An upper bound below zero supports the joint statement that every mean difference is negative. For within-family contrasts, Holm adjustment is applied across the five fully coupled regimes separately for each prespecified contrast. The eight parameter profiles form a separate family, and each preview-scale contrast forms its own five-regime family. A directional cell is labeled supported only if its mean favors the first-named method, its paired interval excludes zero, and its family-adjusted sign-test value is below 0.05. We report seed-level differences and leave-one-seed-out ranges for the prespecified centered full-action versus standard DQN contrast. Its conservative sensitivity analysis resamples the two sets of 20 seed means independently and ignores matching. Test traces and bootstrap draws are never treated as independent model replicates.

3.4.5. Prespecified sensitivity and training diagnostics

Without retraining or model selection, we repeated the nominal evaluation under eight profiles. They covered the reference configuration, compute demand $\pm 20\%$, raw data volume $\pm 20\%$, energy use $+20\%$, heat load $+20\%$, and link capacity -20% . Each profile retained the same 20 models and held-out trace namespace as the primary experiment. Holm adjustment covered all eight profiles. Training logs record every episode for all 20 seeds. Curves report a causal 25-episode moving mean and the seed-level mean $\pm$ standard deviation (SD) for total cost, energy use, latency, and delayed penalty. These

diagnostics describe optimization trajectories; confirmatory claims remain based on held-out traces.

3.4.6. Preview-model misspecification test

The main experiment assumes that the analytical preview matches the simulator transition. To test this dependence, we trained two additional centered full-action model sets without changing the factual environment or deployment evaluation. One set scaled preview one-step costs and action-induced resource changes by 0.85, while the other used 1.15. The preview cost comprised the immediate cost and modeled penalty. Resource predictions were clipped to the same physical bounds after scaling around the current resource state. Each condition used the same 20 seeds, training budget, network, replay, and held-out trace namespace as the main experiment. These symmetric stress-test scales were fixed in advance and were not used for model selection.

3.4.7. Reproducible result entry point

The core result entry point is `python -m dqn_family_satellite_ground.reviewer_experiments all -seeds 800-819 -resume`. Component and Double DQN comparisons use `python -m dqn_family_satellite_ground.extended_revision -seeds 800-819 -stage all -resume`. The first command trains or loads 60 primary learned-policy checkpoints and 40 preview-error checkpoints. The second reuses the primary DQN checkpoints and trains 80 component and Double DQN checkpoints. It then evaluates six DQN objectives, performs the independent-seed bootstrap, and evaluates four planning depths on 400 unique traces.

The command `python -m dqn_family_satellite_ground.extended_reporting` recomputes the 30 paired DQN-MPC-4 contrasts, joint maximum statistics, LaTeX tables, and artifact hashes from locked raw comma-separated value (CSV) files. The command `python -m dqn_family_satellite_ground.transition_audit` writes a fixed-seed, all-action example. In this example, preview and executed costs and all six resources agree coordinate by coordinate. Unit tests repeat this comparison at both clipping boundaries and on the terminal step. They also verify that time and link phase occupy the three added state coordinates.

Experiments used Python 3.13.9, NumPy 2.4.1, pandas 3.0.3, and PyTorch 2.8.0 with CUDA 12.9. Hardware comprised an NVIDIA GeForce RTX 5070 Ti running Windows 11. Planning wall time was measured on a 12th Gen Intel Core i9-12900K with 24 logical processors after one untimed warm-up episode. Reported values are means over all timed decisions. Machine-readable metadata record the 24-dimensional state definition, resolved variants, seeds, budgets, randomization blocks, model-information counts, result files, and runtime environment. Table and figure scripts consume these files directly.

4. Results

4.1. Direct paired inference supports a DQN-family advantage under full coupling

We first tested the family-level claim against the strongest implemented non-DQN comparator. MPC-4 had the lowest non-DQN mean in every fully coupled regime. All 30 paired objective-regime comparisons produced lower mean costs for DQN, and every individual 95% bootstrap interval excluded zero.

Table 4. Episode total cost under the internally consistent synthetic generator. For learned policies, standard deviation (SD) is calculated across 20 independently trained model-seed blocks after within-seed trace averaging. For deterministic non-learned policies, SD is calculated across 20 disjoint trace namespaces and does not represent training variability. These quantities describe different sources of variation and are not directly comparable as inferential uncertainty. All learned policies were trained only at nominal $c = 1$; static and reduced-coupling rows are transfers without retraining. MPC-4 is exact only for its stated four-step perfect-forecast objective. Lower values are better.

Regime	Argmin	MPC-4	Bandit	Threshold	DQN	Centered FA
Static ($c = 0$)	24.55 $\pm$ 0.21	24.55 $\pm$ 0.21	24.78 $\pm$ 0.33	24.74 $\pm$ 0.25	26.85 $\pm$ 0.99	26.66 $\pm$ 0.64
Reduced ($c = 0.5$)	26.65 $\pm$ 0.63	25.70 $\pm$ 0.26	26.39 $\pm$ 1.05	26.43 $\pm$ 0.38	27.81 $\pm$ 0.68	27.61 $\pm$ 0.56
Nominal	193.42 $\pm$ 8.33	90.83 $\pm$ 1.71	101.00 $\pm$ 4.94	188.08 $\pm$ 6.41	78.33 $\pm$ 1.88	77.66 $\pm$ 1.84
Burst	241.20 $\pm$ 6.84	108.35 $\pm$ 1.99	136.55 $\pm$ 11.20	238.74 $\pm$ 6.56	97.36 $\pm$ 2.01	96.88 $\pm$ 1.94
Link-limited	272.02 $\pm$ 7.98	115.73 $\pm$ 1.92	139.75 $\pm$ 7.36	227.25 $\pm$ 8.36	102.26 $\pm$ 1.75	101.80 $\pm$ 1.80
Energy-limited	225.31 $\pm$ 7.73	126.00 $\pm$ 1.58	133.26 $\pm$ 4.20	246.46 $\pm$ 6.45	119.67 $\pm$ 2.14	119.00 $\pm$ 1.65
Thermal-stress	247.58 $\pm$ 6.25	109.20 $\pm$ 2.01	137.74 $\pm$ 10.79	241.80 $\pm$ 6.26	97.55 $\pm$ 2.03	97.10 $\pm$ 1.93




Figure 4. Performance across seven regimes under internally consistent synthetic task generation. For learned policies, error bars show SD across 20 model-seed blocks after within-seed trace averaging. Deterministic comparators show variation across trace namespaces, so their error bars are descriptive and not inferentially interchangeable with those of learned methods. All learned models were trained at nominal $c = 1$; the $c < 1$ rows are transfers without retraining. The ordinate differs among panels, so panel heights should not be compared directly.

Table 5. Direct joint inference for the DQN family relative to MPC-4. Differences are DQN minus MPC-4, so negative values favor DQN. Within each regime, the table reports the largest, least-favorable mean among six DQN-objective contrasts; the final row takes the maximum across all 30 objective–regime cells. Intervals are percentile intervals for the corresponding maximum from 10,000 joint bootstrap resamples. Simultaneous one-sided 95% upper bounds use the maximum centered bootstrap deviation. The independent unit is the model-seed block ($n = 20$), after averaging 20 shared held-out traces within each block.

Scope	Least-favorable DQN objective	Maximum Δ	95% interval for maximum	Simultaneous 95% upper bound
Nominal	Double DQN	-12.43	[-12.92, -11.81]	-11.82
Burst	Double DQN	-10.74	[-11.26, -10.13]	-10.17
Link-limited	Double DQN	-13.15	[-13.65, -12.58]	-12.62
Energy-limited	Double DQN	-5.92	[-6.66, -4.94]	-5.04
Thermal-stress	Double DQN	-11.40	[-11.94, -10.77]	-10.84
All five regimes	Double DQN (Energy-limited)	-5.92	[-6.66, -4.94]	-5.03

The joint seed-block analysis supported the same conclusion after multiplicity correction across all 30 cells. Every seed block favored DQN, and the largest Holm-adjusted sign-test value was $5.722\text{e-}05$. The least-favorable cell was Double DQN under energy limitation. Its mean difference was -5.92, with a simultaneous 95% upper bound of -5.03 (Table 5; Fig. 1). Nominal DQN-family means spanned 77.66–78.40, compared with 90.83 for MPC-4. Together, these analyses support a DQN-family advantage as the central benchmark finding.

4.2. Target design produces smaller, scenario-dependent differences within the DQN family

Table 6. Episode total cost across six DQN objectives in the five fully coupled regimes. Values are mean $\pm$ SD across 20 independently trained model-seed blocks after within-seed averaging of 20 shared held-out traces. Lower values are better, and the leading objective varies across regimes.

Training objective	Nominal	Burst	Link-limited	Energy-limited	Thermal-stress
Standard DQN	78.33 $\pm$ 1.88	97.36 $\pm$ 2.01	102.26 $\pm$ 1.75	119.67 $\pm$ 2.14	97.55 $\pm$ 2.03
Full-action Q auxiliary	78.14 $\pm$ 2.23	97.21 $\pm$ 2.25	102.17 $\pm$ 2.09	119.78 $\pm$ 2.38	97.41 $\pm$ 2.25
Immediate advantage auxiliary	77.84 $\pm$ 1.82	97.22 $\pm$ 2.05	101.95 $\pm$ 1.88	119.17 $\pm$ 2.12	97.40 $\pm$ 2.04
Centered full-action	77.66 $\pm$ 1.84	96.88 $\pm$ 1.94	101.80 $\pm$ 1.80	119.00 $\pm$ 1.65	97.10 $\pm$ 1.93
Double DQN	78.40 $\pm$ 1.98	97.61 $\pm$ 2.15	102.59 $\pm$ 2.08	120.08 $\pm$ 2.41	97.80 $\pm$ 2.12
Double + centered full-action	77.73 $\pm$ 1.99	96.90 $\pm$ 2.02	101.91 $\pm$ 1.88	119.07 $\pm$ 1.77	97.13 $\pm$ 2.01




Figure 5. DQN-family performance across fully coupled regimes. Points show model-seed means and error bars show ± 1 SD across 20 independently trained model seeds. The figure shows overlapping distributions and scenario-dependent ordering across objectives.

Table 7. Nominal fixed-coefficient component and backbone comparisons. All models use 20 pre-specified model-seed blocks, identical real-interaction budgets, and the same held-out traces. Full-action Q and immediate advantage receive the same three modeled action outcomes as the centered full-action variant. The shared $\lambda = 0.3$ was calibrated for the centered objective and not tuned separately for the other auxiliary losses, so these rows do not isolate centering or bootstrapping. Holm adjustment covers the five fully coupled regimes separately for each displayed contrast. “Supported” requires a favorable mean, a paired interval below zero, and Holm $p < 0.05$.

Training objective	Mean cost $\pm$ SD	Matched reference	Paired difference [95% interval]	Holm p	Supported
Standard DQN	78.33 $\pm$ 1.88	—	—	—	—
Full-action Q auxiliary	78.14 $\pm$ 2.23	Standard DQN	-0.19 [-0.56, 0.28]	0.05909	No
Immediate advantage auxiliary	77.84 $\pm$ 1.82	Standard DQN	-0.49 [-0.90, -0.12]	1	No
Centered full-action	77.66 $\pm$ 1.84	Standard DQN	-0.67 [-0.99, -0.39]	0.01288	Yes
Double DQN	78.40 $\pm$ 1.98	—	—	—	—
Double + centered full-action	77.73 $\pm$ 1.99	Double DQN	-0.67 [-1.18, -0.24]	0.03545	Yes

Table 8. Paired inference for centered full-action minus standard DQN across the five fully coupled regimes. Negative differences favor centered full-action supervision. Each comparison uses 20 model-seed blocks after averaging 20 shared held-out traces. Holm adjustment covers the five prespecified regimes, and LOO denotes the leave-one-seed-out mean range.

Regime	Mean difference	95% interval	Negative seeds	Holm p	LOO range
Nominal	-0.67	[-0.99, -0.39]	17/20	0.01288	[-0.71, -0.55]
Burst	-0.48	[-0.73, -0.28]	17/20	0.01288	[-0.51, -0.37]
Link-limited	-0.46	[-0.77, -0.17]	15/20	0.08278	[-0.52, -0.38]
Energy-limited	-0.68	[-1.17, -0.21]	16/20	0.03545	[-0.79, -0.51]
Thermal-stress	-0.45	[-0.75, -0.22]	15/20	0.08278	[-0.48, -0.33]

Table 9. Training information and computation accounting. Wall time is mean $\pm$ SD across 20 models. “Modeled action outcomes” counts analytical action evaluations used by the learning objective. Information and computational budgets additionally reflect these modeled action outcomes.

Training objective	Wall time per model (s)	Real interactions	Updates	Modeled action outcomes	Models
Standard DQN	37.4 $\pm$ 10.4	38,400	9,569	0	20
Full-action Q auxiliary	46.0 $\pm$ 1.4	38,400	9,569	115,200	20
Immediate advantage auxiliary	42.9 $\pm$ 1.1	38,400	9,569	115,200	20
Centered full-action	49.0 $\pm$ 3.1	38,400	9,569	115,200	20
Double DQN	36.0 $\pm$ 1.6	38,400	9,569	0	20
Double + centered full-action	53.2 $\pm$ 3.9	38,400	9,569	115,200	20

Across the five fully coupled regimes, the six objectives occupied a narrow performance band (Table 6; Fig. 5). Nominal means ranged from 77.66 to 78.40. The leading objective varied by regime, and the gap to MPC-4 exceeded the differences among target designs.

Centered full-action produced favorable mean differences against standard DQN in all five regimes and met the complete support rule in 3 (Table 8). The remaining component and Double DQN contrasts produced smaller, scenario-dependent changes in ordering (Table 7). All objectives used the same real-interaction budget, whereas the auxiliary objectives additionally consumed modeled action outcomes (Table 9).

4.3. The DQN family remains below deeper exact planning in nominal cost

Table 10. Nominal receding-horizon planning over 400 unique held-out traces. Each planner has perfect knowledge of the next H tasks and exogenous link states and exactly enumerates 3^H sequences before every action. Runtime was measured on the same workstation and includes only planner action selection after environment construction. DQN inference time was not benchmarked in this experiment, so the table is not a matched deployment-latency comparison.

Planning horizon	Enumerated sequences	Mean cost	Time per decision (ms)	Held-out traces
1	3	99.33	0.13	400
2	9	92.80	0.50	400
4	81	90.83	4.82	400
6	729	89.35	42.50	400

Increasing planning depth reduced nominal mean cost from 99.33 at $H = 1$ to 89.35 at $H = 6$. Over the same range, mean decision time increased from 0.13 to 42.50 ms. At $H = 6$, the nominal MPC cost remained above the DQN-family range of 77.66–78.40. Thus, within the tested horizons, deeper exact planning did not close the nominal cost gap.

Table 11. Engineering outcome decomposition for standard DQN relative to MPC-4. Standard DQN is shown as a conventional representative of the six-objective family. Delta columns report standard DQN minus MPC-4; negative values denote lower use or cost. Low-energy steps are soft-threshold exposure counts for decisions with remaining energy below 0.18.

Regime	Δ immediate	Δ penalty	Δ energy	Δ latency (s)	Δ data (Mbit)
Nominal	0.76	-13.26	-0.104	-1.04	-543.9
Burst	0.42	-11.40	-0.098	-0.92	-489.8
Link-limited	-0.57	-12.91	-0.130	-1.46	-388.1
Energy-limited	1.36	-7.68	-0.102	-0.93	-601.5
Thermal-stress	0.32	-11.97	-0.100	-0.97	-482.5

Regime	Low-energy steps DQN/MPC-4	Min. energy DQN/MPC-4	Max. heat DQN/MPC-4
Nominal	23.6/28.5	0.00/0.00	0.192/0.190
Burst	28.8/32.3	0.00/0.00	0.192/0.190
Link-limited	30.2/35.0	0.00/0.00	0.193/0.191
Energy-limited	38.8/41.8	0.00/0.00	0.193/0.190
Thermal-stress	28.8/32.4	0.00/0.00	0.364/0.362

Table 11 provides a representative decomposition of the gap between standard DQN and MPC-4. Relative to MPC-4, standard DQN reduced delayed penalties in every fully coupled regime while changing immediate cost by only -0.57 to 1.36 . It also used less modeled energy, had lower modeled latency, and transmitted less data. Both methods reached zero in the clipped energy state, whereas standard DQN produced slightly higher maximum heat. This resource-use decomposition applies only to standard DQN.

4.4. Representative DQN training diagnostics show empirical stabilization




Figure 6. Representative training diagnostics for standard DQN and the centered full-action DQN variant. Lines are means across 20 model seeds after a causal 25-episode moving average; shaded regions are ± 1 SD. Curves describe the training distribution and are not used as independent test observations or evidence for the whole DQN family.

All four logged quantities declined during approximately the first 220 episodes and then fluctuated around stable ranges (Fig. 6). The late-episode trajectories indicate empirical stabilization for these two representative family members.

4.5. Sensitivity of the within-family comparison to parameter changes

Table 12. Prespecified parameter sensitivity for centered full-action minus standard DQN. Negative differences favor the centered full-action variant; $n = 20$ independent paired model seeds in each profile. Holm adjustment covers the eight displayed profiles.

Parameter profile	Mean difference	95% interval	Improvement	Holm p
Reference	-0.67	[-0.99, -0.39]	0.85%	0.01804
Compute demand -20%	-0.49	[-0.78, -0.22]	0.79%	0.01804
Compute demand +20%	-0.89	[-1.35, -0.53]	0.98%	0.00322
Input data -20%	-1.03	[-1.55, -0.62]	1.66%	0.01804
Input data +20%	-0.44	[-0.69, -0.23]	0.50%	0.01804
Energy use +20%	-0.60	[-0.98, -0.28]	0.57%	0.01804
Heat load +20%	-0.67	[-1.01, -0.40]	0.86%	0.01804
Link capacity -20%	-0.64	[-0.96, -0.35]	0.72%	0.01804




Figure 7. Parameter sensitivity and planning boundary. (a) Mean paired differences between centered full-action and standard DQN with 95% paired bootstrap intervals. (b) Immediate, four-step rolling, and learned scheduling across coupling and stress regimes.

The paired interval for centered full-action minus standard DQN remained below zero in all eight profiles. Relative differences ranged from 0.50% to 1.66%. The largest difference occurred when raw data volume decreased by 20%, whereas the smallest occurred when it increased by 20%. The within-family direction therefore remained stable across the prespecified parameter envelope.

4.6. Preview-scale perturbations bound one model-assisted variant

Table 13. Effect of training-time preview misspecification. Differences are relative to the same standard DQN models evaluated in the factual simulator. Scaling applies only to the centered full-action variant’s training-time preview cost and action-induced resource changes.

Training preview	Regime	Mean difference	95% interval	Negative seeds	Holm <i>p</i>	Supported
Simulator-exact preview	Nominal	-0.67	[-0.99, -0.39]	17/20	0.01288	Yes
	Burst	-0.48	[-0.73, -0.28]	17/20	0.01288	Yes
	Link-limited	-0.46	[-0.77, -0.17]	15/20	0.08278	No
	Energy-limited	-0.68	[-1.17, -0.21]	16/20	0.03545	Yes
	Thermal-stress	-0.45	[-0.75, -0.22]	15/20	0.08278	No
Consequences scaled to 0.85	Nominal	-0.45	[-0.85, -0.12]	18/20	0.002012	Yes
	Burst	-0.41	[-0.72, -0.18]	17/20	0.01031	Yes
	Link-limited	-0.37	[-0.62, -0.13]	13/20	0.5264	No
	Energy-limited	-0.27	[-0.71, 0.14]	13/20	0.5264	No
	Thermal-stress	-0.37	[-0.70, -0.10]	16/20	0.03545	Yes
Consequences scaled to 1.15	Nominal	-0.77	[-1.11, -0.49]	20/20	9.537e-06	Yes
	Burst	-0.48	[-0.74, -0.29]	19/20	0.0001602	Yes
	Link-limited	-0.33	[-0.62, -0.06]	13/20	0.2632	No
	Energy-limited	-0.92	[-1.45, -0.46]	17/20	0.00773	Yes
	Thermal-stress	-0.42	[-0.71, -0.18]	15/20	0.08278	No

Uniform 15% under- and over-scaling retained favorable mean differences between centered full-action and standard DQN (Table 13). The two conditions met the complete support rule in 3 and 3 regimes, respectively. Their largest Holm-adjusted sign-test values were 0.5264 and 0.2632.

5. Discussion

Six DQN objectives retained a common advantage over MPC-4 across changes in supervision, centering, bootstrapping, and Double DQN evaluation. Their narrow internal

spread suggests that the shared sequential value-learning formulation contributed more to performance than the specific target design.

The contextual bandit comparison and planning-depth analysis support a temporal interpretation of this result. The DQN family achieved lower episode costs than immediate-reward learning, while deeper exact enumeration progressively narrowed the cost gap. Together, these observations are consistent with learned values capturing consequences that persist beyond the current allocation.

The benchmark also makes this comparison traceable and reusable. Shared task primitives generate heat, energy, latency, and transmitted volume, while transition audits connect analytical previews to executed outcomes. Scripts link checkpoints and seed-level estimates to tables, figures, and validation hashes. This evidence chain provides a reproducible platform for comparing sequential decision objectives.

These findings complement work on simulated experience, model-based acceleration, and action-gap learning [15,16,21]. The evaluated full-action variants use analytical one-step outcomes as auxiliary labels rather than inserting simulated transitions into replay. However, factual-only DQN controls also shared the family-level advantage. The evidence therefore supports long-horizon value learning but does not attribute the advantage to one model-assisted supervision mechanism.

5.1. Scope and limitations

The scope of these conclusions reflects the use of a synthetic simulator with fixed dimensionless utilities and clipped, soft-constrained resource states; mission-level applicability therefore remains to be assessed. All learned policies were trained at $c = 1$, while the reduced-coupling evaluations jointly vary dynamics and reward scaling. The planning comparison uses perfect forecasts, and a computationally matched evaluation of DQN inference latency remains outside the present scope. Within-family comparisons use an auxiliary weight calibrated for the centered objective, so the individual contributions of centering and bootstrapping remain to be separated factorially. Preview misspecification was examined using uniform $\pm 15\%$ scaling; broader and structured uncertainty models would further test robustness. These considerations define the scope of the present benchmark and motivate mission-calibrated validation, while the paired comparative result remains supported under the evaluated conditions.

5.2. Future work

Operational validation should replace fixed dimensionless weights with mission-derived utilities and physical resource dynamics. Action masking, constrained reinforcement learning, and terminal resource conditions can enforce feasibility. Telemetry and hardware-in-the-loop traces can then test whether the observed ordering persists. Multi-satellite, multi-ground-station, and variable-contact topologies can assess scalability.

Mechanism studies should tune auxiliary weights by objective, normalize gradient contributions, and factorially separate centering, Bellman bootstrapping, and discounting. They should also vary transition carryover and delayed-penalty scaling independently. Regime-specific retraining can complement the current $c = 1$ transfer tests. Structural, state-dependent, and action-specific preview errors can reveal which components sustain the observed pattern. Family-wide action-mixture analyses can provide complementary evidence.

A broader comparator frontier should include scalable tree search, dynamic programming, mixed-integer optimization, and satellite-specific schedulers. These methods should be evaluated under matched information and computation budgets. DQN inference latency and energy should be measured on the same deployment path as plan-

ning. Such studies can assess whether the benchmark advantage transfers to operational scheduling.

6. Conclusions

The results support a DQN-family advantage in a reproducible, resource-coupled satellite-ground scheduling benchmark with soft constraints. All 30 paired objective-regime contrasts favored DQN over MPC-4, with a family-wide simultaneous 95% upper bound of -5.03. Differences among DQN objectives were smaller and scenario dependent, reinforcing the consistency of the family-level result. The benchmark provides a reproducible basis for mission-calibrated, hard-constrained, and computationally matched validation.

Author Contributions: Conceptualization, B.Y.; methodology, B.Y.,L.Z.,D.L.; software, B.Y.,D.L.; validation, B.Y.; formal analysis, B.Y.; investigation, B.Y.; resources, B.Y.,D.L.; data curation, B.Y.; writing—original draft preparation, B.Y.; writing—review and editing, B.Y.; visualization, B.Y.; supervision, B.C.,J.L.; project administration, B.C.,J.L.; funding acquisition, B.C.,J.L. All authors have read and agreed to the published version of the manuscript.

Funding: This work was supported in part by the National Key Research and Development Program of China under Grants 2022YFF0503904 and 2022YFD2401202; in part by the Shenzhen Higher Education Institutions Stabilization Support Program under Grant GXWD20220811163556003; in part by the National Natural Science Foundation of China under Grants 62202127 and U23A20280; in part by the Shenzhen Science and Technology Program under Grants JCYJ20241202123731040, ZDCY20250901111401002, ZDCY20250901111500001, and ZDCY20250901105202003; in part by the Guangxi Science and Technology Base and Talent Special Project under Grants Gui Ke AD25069103 and Gui Ke DHKL2413 (Research and Application of Key Technologies for Precise Navigation); in part by the Low-Altitude Airspace Strategic Program Portfolio under Grant Z25306106; and in part by the Guangdong Provincial Key R&D Program under Grant 2026B0909050003.

Data Availability Statement: The versioned reproducibility package contains the simulator, 180 trained checkpoints, training curves, locked raw evaluations, seed-level summaries, and direct DQN-MPC-4 inference. It also provides transition audits, Python and MATLAB figure-generation source code, and LaTeX table generators under `dqn_family_satellite_ground/`, together with the current manuscript PDF. The package covers the six DQN objectives, contextual-bandit comparator, preview-error tests, parameter sensitivity, exact MPC planning-depth analysis, joint bootstrap, and maximum-statistic results reported here. The repository does not include pre-rendered manuscript figures or LaTeX source files. Scripts write regenerated outputs to `reproduced_outputs/`. Machine-readable metadata, a manifest, and SHA-256 checksums record seeds, budgets, package versions, source files, and artifacts. The complete versioned package is publicly available at https://github.com/noob32123/dqn_family_satellite_ground.

Acknowledgments: During the preparation of this manuscript, the authors used OpenAI GPT-5.6-Sol for English-language polishing. The authors reviewed and edited the output and take full responsibility for the content of this publication.

Conflicts of Interest: The authors declare no conflicts of interest.

References

1. Pang, Y.; Zhang, Y.; Wang, Y.; Wei, X.; Chen, B. SOCNet: A Lightweight and Fine-Grained Object Recognition Network for Satellite On-Orbit Computing. *IEEE Transactions on Geoscience and Remote Sensing* **2022**, *60*, 1–13. Art. no. 5632913, <https://doi.org/10.1109/TGRS.2022.3216215>.
2. Pang, Y.; Zhang, Y.; Kong, Q.; Wang, Y.; Chen, B.; Cao, X. SOCDet: A Lightweight and Accurate Oriented Object Detection Network for Satellite On-Orbit Computing. *IEEE Transactions on Geoscience and Remote Sensing* **2023**, *61*, 1–15. Art. no. 5608115, <https://doi.org/10.1109/TGRS.2023.3269642>.

3. Pang, Y.; Zhang, Y.; Wang, Y.; Wei, X.; Chen, B. Exploring Model Compression Limits and Laws: A Pyramid Knowledge Distillation Framework for Satellite-on-Orbit Object Recognition. *IEEE Transactions on Geoscience and Remote Sensing* **2024**, *62*, 1–13. Art. no. 5603313, <https://doi.org/10.1109/TGRS.2023.3348470>. 532–534
4. Pang, Y.; Yao, L.; Luo, Y.; Dong, C.; Kong, Q.; Chen, B. RepSViT: An Efficient Vision Transformer Based on Spiking Neural Networks for Object Recognition in Satellite On-Orbit Remote Sensing Images. *IEEE Transactions on Geoscience and Remote Sensing* **2024**, *62*, 1–16. Art. no. 5610916, <https://doi.org/10.1109/TGRS.2024.3367709>. 535–537
5. Zhang, Y.; Ye, M.; Zhu, G.; Liu, Y.; Guo, P.; Yan, J. FFCA-YOLO for Small Object Detection in Remote Sensing Images. *IEEE Transactions on Geoscience and Remote Sensing* **2024**, *62*, 1–15. Art. no. 5611215, <https://doi.org/10.1109/TGRS.2024.3363057>. 538–539
6. Casalicchio, E.; Magliarisi, D. Decentralized Task Scheduling in Satellite Edge Computing. In Proceedings of the 2024 9th International Conference on Fog and Mobile Edge Computing (FMEC). IEEE, 2024, pp. 154–161. <https://doi.org/10.1109/FMEC62297.2024.10710288>. 540–542
7. Kim, J.; Ahn, J.; Choi, H.L.; Cho, D.H. Task scheduling of agile satellites with transition time and stereoscopic imaging constraints. *Journal of Aerospace Information Systems* **2020**, *17*, 285–293. <https://doi.org/10.2514/1.1010775>. 543–544
8. Liu, W.; Lai, Z.; Wu, Q.; Li, H.; Zhang, Q.; Li, Z.; Li, Y.; Liu, J. In-Orbit Processing or Not? Sunlight-Aware Task Scheduling for Energy-Efficient Space Edge Computing Networks. In Proceedings of the IEEE INFOCOM 2024-IEEE Conference on Computer Communications. IEEE, 2024, pp. 881–890. <https://doi.org/10.1109/INFOCOM52122.2024.10621268>. 545–547
9. Jiang, Q.; Wang, H.; Kong, Q.; Zhang, Y.; Chen, B. On-orbit remote sensing image processing complex task scheduling model based on heterogeneous multiprocessor. *IEEE Transactions on Geoscience and Remote Sensing* **2023**, *61*, 1–18. Art. no. 1001718, <https://doi.org/10.1109/TGRS.2023.3327279>. 548–550
10. Jiang, Q.; Zheng, L.; Zhou, Y.; Liu, H.; Kong, Q.; Zhang, Y.; Chen, B. Efficient On-Orbit Remote Sensing Imagery Processing via Satellite Edge Computing Resource Scheduling Optimization. *IEEE Transactions on Geoscience and Remote Sensing* **2025**, *63*, 1–19. Art. no. 1000519, <https://doi.org/10.1109/TGRS.2025.3528015>. 551–553
11. Shen, X.; Lu, Z.; Li, L.; Chen, Y.; Li, X.; Wang, J.; Yao, W. Multistrip Stitching Imaging Mission Planning Method for SAR Satellite Regional Mapping Considering Onboard Energy Consumption. *IEEE Transactions on Geoscience and Remote Sensing* **2025**, *63*, 1–13. <https://doi.org/10.1109/TGRS.2024.3513483>. 554–556
12. Yao, W.; Shen, X.; Zhang, G.; Lu, Z.; Wang, J.; Song, Y.; Li, Z. A spiking neural network based proximal policy optimization method for multi-point imaging mission scheduling of earth observation satellite. *Swarm and Evolutionary Computation* **2025**, *94*, 101867. <https://doi.org/10.1016/j.swevo.2025.101867>. 557–559
13. Mnih, V.; Kavukcuoglu, K.; Silver, D.; Rusu, A.A.; Veness, J.; Bellemare, M.G.; Graves, A.; Riedmiller, M.; Fidjeland, A.K.; Ostrovski, G.; et al. Human-level control through deep reinforcement learning. *Nature* **2015**, *518*, 529–533. <https://doi.org/10.1038/nature14236>. 560–562
14. Watkins, C.J.C.H.; Dayan, P. Q-learning. *Machine Learning* **1992**, *8*, 279–292. <https://doi.org/10.1007/BF00992698>. 563
15. Sutton, R.S. Dyna, an integrated architecture for learning, planning, and reacting. *ACM SIGART Bulletin* **1991**, *2*, 160–163. <https://doi.org/10.1145/122344.122377>. 564–565
16. Gu, S.; Lillicrap, T.; Sutskever, I.; Levine, S. Continuous Deep Q-Learning with Model-Based Acceleration. In Proceedings of the Proceedings of the 33rd International Conference on Machine Learning, 2016, Vol. 48, *Proceedings of Machine Learning Research*, pp. 2829–2838. 566–568
17. Forney, A.; Pearl, J.; Bareinboim, E. Counterfactual Data-Fusion for Online Reinforcement Learners. In Proceedings of the Proceedings of the 34th International Conference on Machine Learning, 2017, Vol. 70, *Proceedings of Machine Learning Research*, pp. 1156–1164. 569–571
18. Armengol Urpí, N.; Bagatella, M.; Vlastelica, M.; Martius, G. Causal Action Influence Aware Counterfactual Data Augmentation. In Proceedings of the Proceedings of the 41st International Conference on Machine Learning, 2024, Vol. 235, *Proceedings of Machine Learning Research*, pp. 1709–1729. 572–574
19. Mesnard, T.; Weber, T.; Viola, F.; Thakoor, S.; Saade, A.; Harutyunyan, A.; Dabney, W.; Stepleton, T.S.; Heess, N.; Guez, A.; et al. Counterfactual Credit Assignment in Model-Free Reinforcement Learning. In Proceedings of the Proceedings of the 38th International Conference on Machine Learning, 2021, Vol. 139, *Proceedings of Machine Learning Research*, pp. 7654–7664. 575–577
20. Kaya, E.C.; Ghasemi, M.; Hashemi, A. Joint MDPs and Reinforcement Learning in Coupled-Dynamics Environments. In Proceedings of the Proceedings of the 42nd Conference on Uncertainty in Artificial Intelligence, 2026, Vol. 337, *Proceedings of Machine Learning Research*, pp. 2877–2893. 578–580
21. Bellemare, M.G.; Ostrovski, G.; Guez, A.; Thomas, P.S.; Munos, R. Increasing the Action Gap: New Operators for Reinforcement Learning. In Proceedings of the Proceedings of the Thirtieth AAAI Conference on Artificial Intelligence, 2016, Vol. 30, pp. 1476–1483. <https://doi.org/10.1609/aaai.v30i1.10303>. 581–583
22. NASA Small Spacecraft Systems Virtual Institute. State-of-the-Art of Small Spacecraft Technology: Small Spacecraft Avionics. <https://www.nasa.gov/smallsat-institute/sst-soa/small-spacecraft-avionics/>, 2026. Accessed 26 August 2026. 584–585

23. NASA Small Spacecraft Systems Virtual Institute. State-of-the-Art of Small Spacecraft Technology: Ground Data Systems and Mission Operations. <https://www.nasa.gov/smallsat-institute/sst-soa/ground-data-systems-and-mission-operations/>, 2026. Accessed 26 August 2026.
24. Consultative Committee for Space Data Systems. Image Data Compression. Technical Report CCSDS 122.0-B-2, CCSDS, 2017.
25. van Hasselt, H.; Guez, A.; Silver, D. Deep Reinforcement Learning with Double Q-Learning. In Proceedings of the Proceedings of the AAAI Conference on Artificial Intelligence, 2016, Vol. 30, pp. 2094–2100. <https://doi.org/10.1609/aaai.v30i1.10295>.

Disclaimer/Publisher’s Note: The statements, opinions and data contained in all publications are solely those of the individual author(s) and contributor(s) and not of MDPI and/or the editor(s). MDPI and/or the editor(s) disclaim responsibility for any injury to people or property resulting from any ideas, methods, instructions or products referred to in the content.